Lichen

L2s In CHannels ENforced by bitcoin

waxwing

Draft, October 2026

Abstract

A *contract* with a hub lets a user hold the hub's tokens on a layer 2 and withdraw them as bitcoin without trusting the hub for honesty or liveness. The L2 tokens are collateralised by the bitcoin in the user-hub multisig output[1], and if the hub refuses a withdrawal, or cannot answer because it is offline, it is enforced on Bitcoin. The only assumptions are those of the layer 2 itself, kept apart from the hub. No change to Bitcoin's consensus is required. A withdrawal is a claim that a program, the withdrawal statement, ran on some input and produced an output: the statement is proved with RISC Zero and wrapped in Groth16, and a disputed claim is settled on Bitcoin by BitVMX [5], whose interactive search over the execution of a RISC-V program, here a Groth16 verifier, is the technical core of the design. Withdrawals are executed cooperatively in the channel — with overwhelming likeliness because the statement is deterministically verifiable in milliseconds — or, otherwise, a refused one is played as BitVMX's on-chain search, pre-signed once per contract. The trustlessness gain is paid for with liquidity (with receiving ability only contingent on offered "inbound" from the hub, as in Lightning). The hub (but not user) must also stay live during normal operation, since an offline hub loses disputes by timeout. The design is built and measured on regtest: a real proof of the withdrawal statement is accepted by BitVMX's unpatched Groth16 verifier in 478 million steps, implying about 22–25 transactions of 8-ary search on chain in the dispute's worst case.

Contents

1 Introduction **2**
1.1 Obvious objections to Lichen and why they're mostly wrong 2
1.2 A note on the delta to LN-GAP 3

2 Contracts with the hub **4**

3 How a withdrawal from the L2 is proved valid **5**
3.1 The statement proved 5
3.2 From a program to a verifiable claim 6

4 Building blocks **6**
4.1 Values that Script can read: one-time signatures 6

[1] This can be embedded into a full Lightning channel, or a single isolated contract; both work.

4.2 BitVMX's search 7

5 The contract logic **8**
5.1 Three ways out 8
5.2 The dispute chain 8
5.3 Outcomes 10

6 Finality from Bitcoin **12**

7 Watchtowers **13**

8 Scaling with multiple hubs **14**

9 Scenarios in which a fault occurs **16**

10 Costs **17**

11 Positioning **18**

12 Open questions **20**

13 Implementation **21**

1. Introduction

1.1. Obvious objections to Lichen and why they're mostly wrong

One central hub is dangerous. If the hub disappears, the user gets the money. The hub cannot Sybil attack users by printing tokens fictitiously; the honest users' money is all backed. The only remaining risks are inherent to the L2 itself (is it a validity rollup or a quorum? Does it inherit finality from Bitcoin or is finality weaker? Etc.)

Bridges usually mean you can only withdraw a fixed size utxo, meaning operators and relying on swaps. Surely that applies here? No. The withdrawal amount is encoded in bits in the bitcoin transaction witness, in dispute. This allows us to withdraw arbitrary amounts via multiple outputs of certain denominations [5.3].

Don't these constructions usually require a 1 of n honesty assumption from a signing committee? Not here: we avoid this by the LN-GAP [1] trick: the contract is a *bilateral* contract between the hub and each user; the presigning happens between them, just as in Lightning.

But if a bunch of users open channels with a hub, they can just .. pay each other through the channels? A correct critique, partly, if the L2 ledger doesn't do anything but plain payments. Even then, you would avoid receiver liveness being needed, avoid routing and other minor things. But the real answer is: the L2 can offer more functionality. It can use zcash style shielded notes, allowing very high privacy while staying in the L2. It can also (in a different "lichen" on a different hub) offer rich smart contracting, etc.

But this an IOU/ hub's token, so it's not actually bitcoin. Incorrect with a crucial caveat. See the first point: every sat is redeemable via the user-hub contracts, *without* the hub's permission, one-to-one, so these are actually bitcoin *but only as long as you have liquidity to receive them.* A new user on L2 without a contract with the hub, or a user holding a balance on L2 far in excess of their inbound liquidity in their own contract, is holding *credit*, and has no such guarantee. But that's their choice.

But still, a single central hub feels like bad design. Why not multiple hubs? Having multiple hubs *will* make sense — see Section 8, but done naively, it is broken. It's because of what these L2 tokens are. They are fungible, enforceable claims on the hub's money; they are minted by the hub itself, just on its own say-so. If you had hub 1 and hub 2, those claims are not transferrable trustlessly to the hubs themselves. If hub 2 prints a billion sats out of nothing and users accept it on the L2 they can redeem it with hub1, leaving hub1 out of pocket. From this line of thinking we see an interesting comparison with ecash servers: the hub can print as much as it likes, but it cannot steal/"rug".

So what's the catch? The negatives are twofold: one is liquidity, and this is easy to understand as the analog of the inbound liquidity pain point in Lightning, because it's basically the same thing. Such a hub *can* give no inbound liquidity to a user: they get tokens on L2 to spend, but they cannot (safely) receive. If they receive, it's just credit that someone else will hopefully buy from them. But the design does not assume that, of course, it allows receiving up to a "cap": the user negotiates that extra potential receiving amount (inbound liquidity) when it sets up the contract with the hub. The second negative is very concrete, and is a design decision rather than an absolute requirement: we use BitVMX's bisection (in fact 8-ary not binary) to enforce that the ZKP is valid, on-chain. In the worst case this is 20-25 transactions and may require watchtower design to be easy enough for users to fully rely on. In a correctly built implementation this is exceptionally rare because the incorrect party (trivially established) is guaranteed to pay all the fees (up to its reserve), and lose money.

1.2. A note on the delta to LN-GAP

LN-GAP[1] puts contracts governed by arbitrary programs into Lightning [8] channels: a contract's outcome is decided by a game whose moves are signed off-chain, with pre-signed transactions that let either party take a dispute to Bitcoin [1]. Lichen mainly differs in trying to achieve trustlessness at the expense of LN-GAP's O(1) onchain resolution; a tradeoff that fits well with the intent of creating *an L2 from a bunch of channels with a hub*:

	LN-GAP	Lichen
venue	trusted: a majority of members attests moves	none
a dispute is decided by	the venue's attestations	the chain: BitVMX's search
suited to	interactive games where time has value	withdrawals from a hub's layer 2
on-chain cost of a dispute	constant	steps on chain log(program length)

Why Lichen needs no venue. What if the "game" the two sides are playing is a fixed, deterministic adjudication of truth? Timing doesn't really matter as much, just ensured completion. Secondly, the actual game playing is going to be a rare back-out clause: unlike say, chess, in which the winner is not predictable at the start, so that the game is the point. If the "game", as here, only exists to prevent cheating and can only play out in one way, it's more important to achieve trustlessness, even at cost of a lot more onchain execution, because that cost won't usually need to be paid and the value of real trustlessness is immense.

2. Contracts with the hub

Alice, the user, has a contract with a **Hub**, instantiated as a 2 of 2 output in a channel between them, or on its own (Section 5), with these terms:

D	Alice's deposit: she puts D bitcoin in and receives D of the hub's tokens on the L2
$V_{\max}$	the most she can withdraw: her D and the hub's liquidity $V_{\max} - D$
r_A, r_H	the *contract reserves*, Alice's and the hub's: returned in every honest outcome, forfeited to the other side by losing a dispute (Section 5.3); the contract holds $V_{\max} + r_A + r_H$
b	the amount she withdraws, by returning b tokens to the hub on the L2 with the memo c
c	the contract's identifier, its memo
T_{close}	the contract's end: with no withdrawal by then, the hub takes it

Every token is a liability of the hub, collateralised by the contract of whoever holds it, within that holder's cap; within $V_{\max}$ and before T_{close}, Alice trusts the hub for nothing, not even to be online: a withdrawal it does not answer is paid by timeout. The hub's provided liquidity is the sum of its users' $V_{\max} - D + r_H$.

Opening is a swap. Alice's D goes into the contract and her D tokens onto the L2, and neither may come without the other: a hub that took the funding and never minted would leave her nothing to return, and the default at T_{close} would hand it her D; a hub that minted first could be left with tokens backed by no contract. Alice first pays a deposit output, the hub's key and hers together, or hers alone after a timeout. The hub then mints D to the key $P = H(A \,\|\, T) \cdot A + T$, where A is Alice's L2 key and $T = t \cdot G$ a point of the funding transaction's own (the coefficient keeps Alice from choosing A against T), and Alice, once she sees that mint sealed, signs the funding's deposit input as an adaptor pre-signature locked to T. The hub can spend the deposit only by completing it, which publishes t and so unlocks her tokens. The hub's side needs no refund: if it never broadcasts, Alice refunds the deposit, and the locked tokens, which cost nothing to print, stay where they are. Its one risk is timing: t is public once the funding is in the mempool, so it broadcasts only well before the deposit's refund matures.

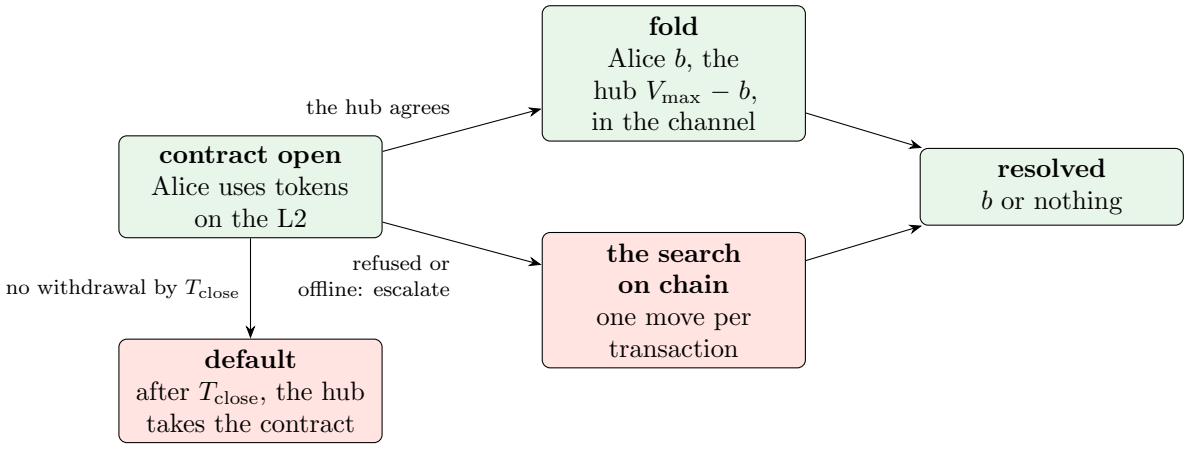


Figure 1: The life of a contract. Green states stay off-chain or settle cooperatively; red states are on chain.

3. How a withdrawal from the L2 is proved valid

3.1. The statement proved

For brevity we have not yet given a description of what the L2 looks like, but its main features are here, to explain the proof structure.

The L2 is a ledger of notes, with state updates signed by a sequencer[2]. A note is an owner, a value, a memo and a random ρ, which makes every note distinct (two equal payments to the same owner are still two notes); its commitment is a SHA-256 hash of these, and the commitments form an append-only Merkle tree. A spent note is not removed from the tree, which would change every other note's path, but recorded by its *nullifier*, a hash of its commitment, in a set the sequencer keeps: spending a note twice repeats its nullifier. So a note's path under a root stays valid however the ledger moves on. This description may confuse a little, but we're trying to avoid adding the complexities of shielding a la Zcash designs *now*, for the proof of concept phase; but we intend to add something in that style: in that case, ρ also hides the note's contents, and the nullifier is derived from a secret key so that nobody can link it to its note. A state root R hashes the height, the tree's root and the nullifier set's root, and the L2's sequencer signs it (BIP340). The statement therefore is:

> "a note owned by the hub, of value b and memo c, is in the commitment tree of a state root the sequencer signed"

Its public output, the *journal*, is $b \,\|\, c$. The program created to generate these proofs has what's called an "image id", which includes both sequencer and hub keys (remember, only the hub can mint L2 tokens).

[2]We'll talk more about achieving finality via Bitcoin or otherwise later [6].

3.2. From a program to a verifiable claim

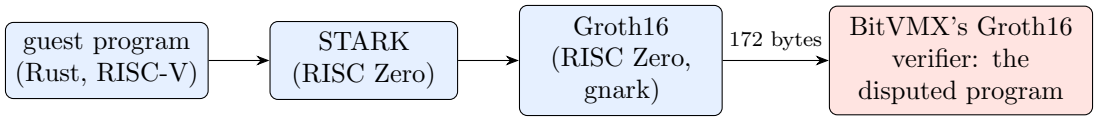


Figure 2: The withdrawal statement runs in RISC Zero's zkVM; its proof is compressed to a Groth16 proof; a dispute is over the execution of a fixed RISC-V program that verifies it [12, 5]. A companion note explains each stage [2].

So the disputed program is always the same: the Groth16 verifier that FairgateLabs compiled for BitVMX [6], unpatched. Its input is 172 bytes, read as 43 four-byte words, each signed by Alice with its own one-time key when she makes her claim:

words	content	used by the contract
0	the journal's length (2)	checked
1–8	the image id	checked against the pinned statement
9–40	the proof's three points[3], compressed	by the search only
41	b	read by the payout
42	c	checked against the contract

The checks are the hub's leaves on the contract: Alice's signature on a word with the wrong value gives the hub the contract. The image-id check is essential: without it, Alice could submit a valid proof of a *different* program, which the verifier would accept.

4. Building blocks

More detail on the contract *construction*. We're thinking of a single output (utxo), 2-of-2 controlled by Alice and the Hub, which is taproot and so can have several spending conditions encoded in leaves of the taproot tree. Apart from the normal segwit trick of pre-signing transaction chains, leaving the witness chooseable after the fact, there are two main building blocks:

4.1. Values that Script can read: one-time signatures

Bitcoin Script can check a signature only on the spending transaction, never on a message of the signer's choosing. A contract that must act on a value a party chooses later (Alice's b, a move in the search) uses Winternitz one-time signatures, which Script verifies by hashing alone, leaving the signed digits on the stack to compare [9, 5]. The keys are exchanged when the contract opens, one per value, each used once; signing two different values under one key is *equivocation*, which a leaf can prove, giving the signer's stake to its counterparty.

[3] A Groth16 ZKP always consists of exactly 3 elliptic curve points, around 100-200 bytes depending on compression serialization.

4.2. BitVMX's search

BitVMX [5] is the main workhorse tech in this design. BitVMX-CPU settles a claim about a RISC-V program's run. The rest of the design details arrange the conditions under which BitVMX's search can serve as the backstop of a withdrawal, using its emulator, search and verifier, with two bug fixes and the authenticated jump list added in a fork [6, 7]. The prover (here Alice) claims that the program, on her input, halts after n steps with a given exit code; each step's write is hashed into a chain, and she commits to the last hash. If the verifier (the hub) disagrees, they search: in each round the prover posts the hashes after $k-1$ evenly spaced steps of the interval in dispute, and the verifier names the segment it disputes. With $k = 8$, ten rounds cover the verifier's 2^{29} steps. They end agreeing on the hash after one step, t, and disagreeing on the hash after the next. The prover then posts her record of step $t+1$: the instruction, the values it reads (and the step that last wrote each) and writes, and the hashes on either side.

Every move is its mover's one-time signatures on the values it posts, one key per field (the claim, each hash, the choice, the record, each step hash), and nothing else. Nothing about the search's state is signed: a rule re-reveals the signatures it reads, which are public once posted, and derives what it needs, as BitVMX's own scripts do. Alice's *proof* of the disputed step, one pre-signed leaf for each class of instruction, runs BitVMX's verification script for the instruction on her record and checks the step hash. The hub's *disproves* are BitVMX's challenges: the agreed step must be the hub's choices concatenated, the step hashes must be those Alice posted for steps t and $t+1$ in the search (an *equivocation* challenge for each round and index), the record must agree with the program (the opcode at its pc, the entry point, the pc the previous step wrote, the memory sections) and its never-written reads with the input or the initial data, and a claimed success must end in a success exit (`halt_exit`).

Reads. The step hash covers what a step writes, not what it reads: a read's value is whatever was last written to its address. So Alice could claim a false value for a read, naming a plausible step as its last write, and execute the instruction faithfully on it: her chain would diverge at that step, the search would find it, and her proof would pass. BitVMX catches this with a second search for the read's last write, as long again as the first. Here the trace is hashed instead as an *authenticated jump list* (S. D. Lerner [13]):

$$h(i) = \mathrm{BLAKE3}\big(h(i-1) \,\|\, h(\mathit{jump}(i)) \,\|\, \mathit{write}_i \,\|\, i\big),$$

where $\mathit{jump}(i)$ is an earlier step chosen by the step number alone, so that any step reaches any earlier one in a few times $\log_2 n$ hops, each to the previous step or to the long jump (at most 86 for 2^{29} steps, about 52 on average). A genuine hash commits to its step number, so its back-links are the genuine hashes of the steps it names. To show a false read, the hub posts the hashes along a path from the agreed $h(t)$ back to the last write, and the write's preimage; its verdict wins if the write there contradicts Alice's read (another value at the step she names, or a later write to her address). Alice, holding the true history up to t, may challenge any hop whose hash she finds false; the hub must then reveal that step's preimage, which must hash to the posted hash and link back to the next, so a forged path loses at its first false hop. The same machinery settles one more lie a jump list makes possible: Alice's disputed step commits to the hash of its long jump, and a false one would let a valid write prove on a false chain; the hub's path to $\mathit{jump}(t+1)$ shows the true hash. This is one round where BitVMX needs a second search: about three more transactions (Section 5.2).

5. The contract logic

5.1. Three ways out

A contract ends in one of three ways; its output, in each commitment, carries two leaves for them besides Lightning's revocation:

- the *fold*: a cooperative update of the channel (Alice $b + r_A$, the hub $V_{\max} - b + r_H$), the normal case, nothing on chain;
- `default`: after T_{close}, a transaction pre-signed at the contract's opening pays the hub $V_{\max} + r_H$ and returns Alice's reserve r_A. This is the liveness requirement on Alice: turn up at the end, i.e. it's a very low bar requirement; whereas the hub must always be ready to answer disputes.
- `escalate`: Alice's claim goes on chain as the first move of BitVMX's search, with her signed input words, and the search continues on the *dispute chain*. This will almost never happen, because given Alice's proof, the hub can calculate $\sim$ immediately that he's guaranteed to lose, so just pays out cooperatively (*fold* as above).

In a channel, or on its own? A contract can be a standalone output, a 2-of-2 between Alice and the hub with the same two leaves, opened by one on-chain transaction and folded by another (a cooperative key-path spend paying $b + r_A$ and $V_{\max} - b + r_H$). Putting this output in a Lightning channel is an efficiency layer: in it the fold is a channel update, collateral can be reallocated and contracts renewed without touching the chain, and deposits can arrive as Lightning payments; the price is that every channel update re-signs the contract's pre-signed transactions, for both commitments. The proof of concept uses the standalone form; earlier LN-GAP code demonstrates contracts embedded in channel commitments, with revocation; it could be applied here.

5.2. The dispute chain

Remember that in dispute we do a search over the execution trace of a program, which is actually a Groth16 verifier.

The dispute search alternates Alice's rounds of hashes and the hub's choices: with an 8-ary search over the verifier's 2^{29} steps, ten rounds and 21 moves. On the dispute chain each move is a transaction spending the previous one. Each of the chain's outputs has two pre-signed spends: the *post*, by which the next mover publishes its signed move and creates the next output, and the *timeout*, valid w blocks later, which settles the dispute against the party that failed to move. Each output also carries the other side's rules on the move just made (a choice out of range, Alice's; the final record's challenges, the hub's), and the last step is proved in Script.

A move. A move is its mover's one-time signatures on what it posts, one key per field: Alice's round of seven hashes, about 7.5 kvB; the hub's choice, a byte, about 0.3 kvB; her final record and its hashes, about 6.8 kvB. A key per field keeps the rules within Script's stack: the stack holds at most 1,000 items, the witness counted, and a signed byte costs four (two digits and two chain values), so a rule can re-reveal only a few hundred signed bytes, and re-reveals only the fields it reads.

What the other side must learn from the chain. Each party runs its own emulator and sees nothing of the other's but what is on chain, so everything it needs to play and to disprove must be published there, under the other's signature. Three things are: move 1 signs the claim (the claimed last step and final hash) with round 1's hashes, and the final move signs the disputed step's record; and the hub, to execute the program at all, needs Alice's whole input, of which the escalation carries only the words the payout and the checks read. When those don't cover the input, an `inputs` link follows the escalation with the rest of her signed words (32 for the Groth16 verifier), before the hub must move; if it does not come within the window, the hub's timeout takes the contract.

Pre-signed once. Each step is a 2-of-2 transaction pre-signed when the contract is opened; the move travels in the witness, which a transaction's id does not cover, so every step's id is fixed in advance and the next can be signed against it. Pre-signing binds a transaction to the exact outputs it spends, so a pre-signed chain can be shared only by paths that reach it through the same transaction. Here there is one: `escalate`. So the dispute chain is pre-signed once, linear in the search's length: about two transactions a move, the final step's proofs (one per instruction class, pre-signed so that a proved withdrawal pays b), and the read and jump challenges' set below.

Optional off-chain play. The parties may still play the search off-chain by agreement, exchanging signed moves; if either stops, the other escalates and *replays* the moves already made from the signatures it holds, then play continues on chain. Since each move is signed alone, the escalator could replay one of its own moves other than the one its opponent answered. But it would have signed two values under one key: every dispute chain output carries an `equiv_d` leaf for each key of each move so far, by which the other side shows the two signatures and takes the output at once, ahead of any disprove, which waits δ. So replay is faithful.

Fees. A link of the dispute chain (the escalation, a move) must keep its transaction id, since the next link was signed against it, so nothing can be added to it. It is an ordinary transaction paying a floor of 1 sat/vB of its own out of the contract, so that it relays alone, with an *anchor*, an output only its poster can spend; the poster tops the fee up with a child spending the anchor and a coin of its own. The opponent's next move waits a block after the link confirms (`post` carries `1 OP_CSV`), so nothing but the poster's child can sit on an unconfirmed link: the challenge and response happen on the chain, not in the mempool, and the pinning that TRUC's mempool rules exist to prevent has no handle here. A link may then be as large as relay policy allows (100 kvB), where TRUC caps it at 10 kvB: the 8-ary escalation is 9–12 kvB, a path's chunks up to 93 kvB. The floors come out of the loser's reserve. Each is sized as if the leaf its link spends sat at its tree's greatest depth (which leaves are deepest depends on their hashes, and so on the keys), so the floors are a property of the statement, fixed before any contract opens: on the Groth16 verifier the deepest dispute, a worst-case read, has 195,411 sats of them at 1 sat/vB, half of it the path's combine. A transaction that ends the chain (a timeout, a proof, the default) is spent only by runtime leaves, so its id may change: both signatures on it are `ALL|ANYONECANPAY`, committing to its own input and all its outputs, and the broadcaster adds a coin as an input. The feerate is chosen at broadcast rather than guessed at signing, beyond the links' floors. The prices are the usual ones: a party keeps fee coins (a coin added to a transaction whose outputs are fixed has no change), and the floors are paid whatever the feerate.

The read and jump challenges. A path (Section 4.2) of up to 86 hashes, 1,720 signed bytes, does not fit one input's stack, which holds about 250. So it is posted in chunks, one per input of a single transaction: from the final record's output, the hub's `path_split` makes eight *carriers* (the first holding the value, the rest dust), and `combine_`k spends the first k of them, chunk c in input c (the first chunk holds the hops' count, their types, the path's end and 8 hashes; each other 11). Both are pre-signed, a combine for each k. After the combine, Alice may post her choice of a hop to challenge; after her window, the hub's verdict (`jump`, `read_`r) takes the output with the end's preimage; after her choice, the hub's reveal of that hop does, or Alice's timeout pays her. If the hub splits and never combines, the first carrier's timeout pays her; the hub sweeps the other carriers, dust, at leisure. A read dispute adds three or four transactions to the search.

5.3. Outcomes

The search only ever decides Alice's withdrawal: b if her claim stands, nothing if it fails; the rest of $V_{\max}$ goes to the hub. The reserves follow the dispute: the loser's goes to the winner, whether it lost by a disprove, a failed input-word check, `equiv_b`, or by not answering in time. A disprove, a runtime spend, pays the disprover the whole output: it fires only on a provable lie, so the liar loses.

Paying b: the bit outputs. When Alice wins, the payout must send her b (as built, rounded down to 512 sats: bits below 2^9 have no output), a value she chose when she claimed; but the payout is pre-signed, so its amounts are fixed, and Script cannot compute an amount. One pre-signed transaction per possible b would be too many. Instead, as for numeric DLCs [18], the payout decomposes b in binary: for a unit u and K bits ($2^K u > V_{\max}$), output i holds $u\,2^i$ and has three leaves:

leaf	spent by	condition
Alice's	Alice, after a delay δ	her one-time signature on b (word 41) has bit i set
the hub's	the hub	the same signature has bit i clear
`equiv_b`	the hub	two signatures on different values of b

Since b is under a one-time key, the bits Alice can claim are those of the single value she signed, and their outputs sum to b; the hub takes the rest. For $b = 9$ she takes outputs 0 and 3. Her leaf waits δ so that the hub can show an equivocation first. With $V_{\max} = 0.1$ BTC and $u = 1000$ sat, $K = 14$; each output's spend is about 300 vB, and each party can sweep its outputs in one transaction. The reserves sit in one more output, paid to the winner, with an `equiv_b` leaf of its own. In the cooperative fold none of this appears: the channel simply pays b.

The contract reserve. A reader who knows Lightning will see a familiar question here. Lightning's penalty deters publishing a revoked state only if the cheat has something to lose: a party whose balance has been drained to nothing can publish an old state for free and profit if its counterparty is not watching, which is why channels keep a *channel reserve*. A contract without reserves has the same hole, on both sides, though not where one might first look:

- What Alice risks by a false claim is not her deposit D but her honest outcome, the return b she could actually prove. Once she has spent her tokens on the layer 2 that is nothing,

and a claim of $V_{\max}$ becomes a free option on the hub being offline. It is cheap to hold: she escalates, and if the hub answers she simply stops and loses by timeout what she would not have had anyway.

- The hub makes no claims; it refuses and challenges. Losing the search costs it only fees, since its honest outcome is what it is left with anyway. So it can refuse a valid withdrawal and hope that Alice cannot play the search in time, whatever it committed to the contract.

So the quantity that matters is not what each party put in but what each loses by losing a dispute, and the remedy is a reserve on each side that only a lost dispute moves. Sizing follows the same logic as Lightning's. A probe that succeeds with chance p (the hub is offline) gains up to $V_{\max} - b$ and otherwise costs the reserve, so it is unprofitable when

$$(1-p)\, r_A > p\,(V_{\max} - b), \quad \text{that is, for every } p < \frac{m}{1+m} \text{ with } r_A = m\,(V_{\max} - b).$$

A reserve can therefore deter probing as firmly as one likes, at a price in locked capital: $m = 1\%$, about a channel reserve, deters only a probe with less than a 1% chance; $m = 1$ one with less than even odds; $m = 5$ anything short of near-certainty that the hub has lost its keys. Nobody would lock several times their stake to hold a contract, just as nobody would keep a channel reserve of 500%; the practical choice is where to sit on that curve, and it depends on what the parties know of each other. A floor is clear in any case: every abandoned probe and every refused valid withdrawal should cost the one who started it at least the other side's dispute fees, so that griefing is paid for by the griefer. Part of those fees is known when the contract is signed: the links' floors, which reach 195,411 sats in the deepest dispute on the Groth16 verifier (Section 5.2), so neither reserve should be smaller. lichen's hub takes for its own the larger of its share of the cap and that figure. The first defence stays what it is in Lightning: the hub is live because that is its business, and Alice, the occasional absentee, is the one a reserve protects. The price is capital (Alice brings $D + r_A$, the hub $V_{\max} - D + r_H$) and severity: an honest party who starts a dispute and then disappears loses its reserve as well.

Who it suits. A reserve is locked, not spent: an honest party gets it back in every outcome (the fold, the default, a dispute it wins) and loses it only by losing a dispute, which for an honest Alice means going absent in the middle of one. So nothing stops a small balance with a large reserve; Alice pays only in liquidity. The binding cost falls on the hub. Its reserve must make Alice whole when it refuses her withdrawal and loses, so it must cover her worst-case dispute, however small her balance: about 195 kvB on the Groth16 verifier, which is 195,000 sats in floors at 1 sat/vB, 400,000 at 2 sat/vB, two million at 10. On top of that the hub locks the cap's headroom $V_{\max} - D$. A user holding 50,000 sats therefore costs the hub about as much locked capital as one holding two million, and earns it less. Alice's own reserve is the hub's to demand, and a hub can ask a small user for less at the price of some exposure to probing. The design therefore suits balances from about half a million sats, and comfortably from a few million, where these reserves are a few percent of the stake. Large balances raise no new cost in the contract: the dispute does not grow with b, whose binary payout only adds outputs. What grows is the trust placed in the layer 2's finality, which anchoring its roots on Bitcoin addresses (Section 6). A contract's amounts are 32-bit counts of a unit, so with a unit of one satoshi a contract holds up to about 42.9 BTC.

6. Finality from Bitcoin

With a single, or indeed a distributed, sequencer signing updates to L2 state, without a consensus mechanism like proof of work or proof of stake, it's free to "rewrite history" for transactions on the L2, enabling double spend attacks. The techniques used in ETH L2s to handle this involve the rollup data publication and contract state on the base chain; that is not available here. So in the absence of proposing expensive consensus on L2, it makes the most sense to inherit finality from Bitcoin, as Shielded CSV [19] and Shielded Bitcoin [20] do (the first by publishing each transaction's nullifier on Bitcoin, the second by publishing every transaction there in full), but here using the single use seal concept [15], with one commitment per batch.

If each sealed batch's root is committed in a Bitcoin transaction that spends the previous batch's commitment output, a fork of the layer 2 needs a double spend on Bitcoin. It costs one small transaction per batch, not per layer 2 transaction. A committee of sequencers then needs only one honest member for inclusion (members append batches in turn); forks are excluded by Bitcoin, not by a threshold.

What sealing does not provide is availability: a seal commits to a root, not to the data behind it. This paper does not answer this question; it is a function of exactly how the L2 ledger is constructed (with a Bitcoin publication component, or not). But if data availability is not addressed, the user does have to trust the sequencer in a very concrete sense: if the sequencer refuses to provide the Merkle path for her return note, the user has no way to retrieve their funds in the contract, and loses them. This strongly argues for the above-mentioned "committee of sequencers", because you just need 1 to include your withdrawal and give you its path to be safe[4].

In this scenario, Alice's withdrawal statement must then show that its root is sealed: a light-client/SPV proof inside the guest, a chain of block headers from a checkpoint to the block carrying the seal, with its Merkle path, and some confirmations after it. The checkpoint is a block hash pinned in the contract when it opens, checked as an input word like the memo; since every withdrawal falls before T_{close}, the chain a proof covers is bounded by the contract's lifetime, and a contract shorter than a difficulty period (2016 blocks) crosses at most one retarget, so the guest can bound each header's target by the checkpoint's within the protocol's factor of four.

The cost is surprisingly small: a header is about three SHA-256 compressions, which RISC Zero accelerates, so a thousand headers are a few hundred thousand cycles, against the 446,000 of today's statement (Section 10). The comparison is like for like: the statement already uses RISC Zero's accelerators for its hashes and the sequencer's signature, and the headers' figure assumes the same SHA-256 circuit, 68 cycles a block [16]. So a thousand headers would roughly double the statement, still under a million cycles; hashed in plain software, at a few thousand cycles a compression, they would cost around ten million, some twenty times today's whole statement. The dispute is unchanged. In the cooperative case the hub checks seals with its own node.

Forgery then becomes a matter of mining a fork in Bitcoin itself. In current Bridge designs, like Union [11], a challenger answers the claimant's chain with a heavier one, which the claimant may dispute in turn, in a second BitVMX game with the roles reversed.

[4] A more advanced idea is publishing withdrawals on Bitcoin, requiring L2 to read L1 in a similar way to what is described next

But we don't want to or need to do that; the claimant is always the user and the dispute is settled on one statement.

So any forgery attack is on the side of the user; and such an attack to achieve maximum returns *could* use it to steal from the hub across multiple contracts. So the bound is not one contract's cap but the hub's total headroom exposed to a single forged chain. As long as that stays well below the reward of the blocks a fork must replace, the attack remains un-economic. This is the kind of security you can claim for Bitcoin itself, basically.

7. Watchtowers

A dispute runs for a day or more (about 140 blocks at $w = 12$, as built), and Alice must answer each of the hub's moves within w blocks. A laptop that sleeps or travels will miss one, and a missed move loses the dispute and her reserve. The natural answer is Lightning's: a *watchtower*, a service that plays her side of the search while she is away. Its moves follow from her trace, which it can compute itself from her public input. The difficulty is that a move is not just a value but her commitment to it: a hash from her trace signed with her one-time keys for that move, and which hash she must sign depends on the hub's earlier choices, so the signatures cannot be made in advance.

Giving the watchtower her one-time keys is probably not a viable solution, since you would be trusting the watchtower to not deliberately post invalid moves in the challenge game, thus losing your money. Watchtowers could be bonded and slashed if they post wrong responses to the hub, but that's an ugly solution.

Committing to the trace. The moves can be made keyless. When she escalates, Alice signs a Merkle root over every step hash of her trace; each later move is then an *opening* of that tree, a step hash with its path to the root, which anyone holding the trace can produce and nobody can forge. Any number of watchtowers could then post her moves, and the worst any of them could do is nothing: Lightning's one-of-n model. The cost here is that the opening must be checked when the move is posted, since an unchecked post would let a dishonest watchtower publish garbage first; without `OP_CAT` a Merkle path cannot be hashed natively, and checking a 29-level path in Script, with a hash function built from arithmetic opcodes and split across inputs to fit the stack limit, would make each move hundreds of times larger than today's. With `OP_CAT` it would be cheap.

User-controlled? In Lichen only the user can start a dispute, so her liveness requirement is on demand: from the moment she escalates, for about a day, plus a deadline (T_{close}) she knows from the start. A standalone contract has nothing the hub can broadcast against her in the meantime. That changes the watchtower question. A machine needed for only a day can be her own: a container on a rented VPS, holding her one-time keys, started when the hub refuses a fold and deleted after. It's trusted like her laptop, and no third-party watchtower is needed for honesty. But note: embedding the contract in a Lightning channel, which we have described as a viable variant above a few times, loses this property, because channels always need revocation.

It seems like this 3rd option is the clear winner, though it might involve various finesses around the idea of "doing Lichen in Lightning". It's also slightly concerning that this per-user solution

for watchtower-ing never really took off in Lightning itself (but at least here it is ephemeral and not usually needed . . .).

8. Scaling with multiple hubs

As per most of the second half of this paper, starting above with considerations for watchtowers, all of the rough edges occur around the dispute: from the cost imposed on the user, to the fees paid on the blockchain in the worst case. All of this points to the worst failure mode, which while difficult for an attacker to carry out, would be very bad indeed: mass exit. As we know even in Lightning there is the idea that, if the number of users is high enough, mass publication of old/invalid states could lead to an excessive reliance on prompt confirmation on a blockchain where space is a scarce resource.

Now imagine lichen similarly, but with that much worse feature: a single mega-hub hosting, let's say, 60,000 users. There, since each disputed withdrawal can require 20 or more transactions worst-case, and a few of those transactions are very large, it's clearly the case that a malicious hub can trigger a very large block space demand and, with a lot of difficulty to be fair, *could* steal from many honest users at once depending on the order in which transactions get confirmed. The same attack exists in reverse: if a user-attack generates a ton of Sybil contracts with one hub, they could trigger mass exit with fraudulent exit requests.

It is worth mentioning that the more likely failure mode — that is of hub infrastructure failure leading to offlineness — does not impose nearly the same burden on blockchain space, and is thus less likely to cause harm to users (though it will certainly cause harm to the hub, which does lose money in cases where it loses liveness for meaningful lengths of time).

How many users can exit? Table 1 counts exits by their bytes alone, from the sizes in Section 10. When one side is hostile a different limit binds. To make an honest party miss a move, the attacker must outbid it in every block of that move's window of w blocks; that costs about the same whether one defender is shut out or many, so it pays only when the moves due in one window could not all fit in it anyway: $Ns > w \cdot 1\,\mathrm{MvB}$, for N contracts whose heaviest move in a window is s. This is Flood & Loot [14] in lichen's form. Table 2 gives the largest N for which the attack does not pay assuming the attacker-crafted worst case: every contract's heaviest move is due in the same window.

We can improve this on *some* axes by having multiple hubs. If we had 20 distinct hubs of 3000 users and not 1 hub of 60,000 it would reduce some risk for users in particular. Which leads to the question:

Can you pay users in a different hub? Certainly you can, with some infrastructure. There's more than one way to do it, but for brevity we describe the simplest: if Hub 1 and Hub 2 exist, with their own distinct ledgers, Alice on Hub 1 can pay Bob on Hub 2, in bitcoin, if Hub 2 has a contract on Hub 1 with inbound capacity. She simply pays Hub2's address on Hub 1's ledger, contingent on an atomic swap set up with the payment to Bob on Hub 2's ledger. There could also be routing, but the principle is clear.

Different ledgers? It seems natural in this multi-hub world to allow users to use a Hub whose ledger rules (e.g. does it use shielded transactions or smart contracts or ..) can be distinct and

exit	per user	per block	per day, per month, at a quarter of each block
hub present: a fold	∼0.3 kvB	3,300	120,000; 3.6M
hub gone: escalate, `inputs`, the hub's timeout, the payout	∼23 kvB	43	1,600; 47,000
hub disputes a true claim to the end	123 kvB	8	290; 8,800
the hub disproves a user's read lie	187 kvB	5	190; 5,800

Table 1: Exits by bytes alone: a day is 144 blocks, a month 4,320. The fold was 193 vB on Mutinynet and the absent-hub exit 23,051 vB; the disputes are as measured on regtest (Section 10). With the hub gone, each exit has until its contract's T_{close}, so staggered expiries spread the load.

	a hostile hub, against honest users	hostile users (read lies), against an honest hub
the defender's heaviest move in one window, s	Alice's proof of the last step, 25.1 kvB (her heaviest otherwise: `inputs`, 8.0)	the hub's `combine_8`, up to 93 kvB
contracts safe at $w = 12$ (two hours)	480	130
at $w = 144$ (a day)	5,700	1,550
at $w = 1{,}008$ (a week)	40,000	10,800
the window needed for 3,000 contracts	75 blocks (13 hours)	280 blocks (2 days)
for 60,000	1,510 blocks (10 days)	5,580 blocks (39 days)

Table 2: Contracts a hub can hold before congesting the chain pays, $N \leq w \cdot 1\,\text{MvB}/s$, with whole blocks and every defender's move due in the same window: hostile users choose when to escalate, and a hostile hub can delay each of its replies by up to w to bunch them. The escalation (11.6 kvB) is not counted, since it runs against T_{close} rather than a window.

whose consensus establishing mechanisms are distinct. Alternatively they could all rely on the same finality and proof of validity, in one ledger, as distinct assets.

9. Scenarios in which a fault occurs

scenario	outcome
the hub disappears	Alice escalates; the hub never answers; the dispute chain's timeout pays b and both reserves
the hub refuses a valid claim	the search on chain; Alice proves the last step; b and both reserves
Alice claims a return she never made	the verifier halts with failure; the hub's `halt_exit` disprove; the hub takes everything, both reserves included
Alice props up a false claim by lying about a read, or about her step's long jump	the hub posts a path through its history to the read's last write (or to the long jump); unless she shows a hop of it false, its verdict takes everything
the hub forges a hash in that path	Alice challenges the hop; the hub cannot reveal it; her timeout pays b and both reserves
Alice probes: a false claim, abandoned when the hub answers	she loses by timeout: her reserve goes to the hub
a proof of another program, or another contract's memo	the hub's input-word check takes the contract
Alice signs two values of b	the hub takes the bit outputs with `equiv_b`
Alice disappears until T_{close}	the default; her reserve returned; her tokens remain as credit
the hub takes the funding and never mints, or mints and is never funded	cannot happen: the mint comes first, locked to the funding's adaptor point; broadcasting the funding unlocks it, and without the funding no one can spend it
either party force-closes the channel while the contract is open	nothing moves; the channel is lost
either party forces the search on chain to cause trouble	the outcome is unchanged; each side pays the fees of its own moves, and the loser's reserve covers the winner's
the escalating party replays one of its own moves altered	the other shows two signatures under one key (`equiv_d`) and takes everything
the sequencer censors or forks	the layer 2's assumptions fail; outside the contract's protection
the sequencer signs a root that does not follow from valid transitions	no chained validity proof exists for it, so no withdrawal can use it (not yet built)

Griefing. In a force-close in Lightning, either side can impose on-chain costs on the other by disputing. The difference, here, is the size: a whole bisection search on chain (here 20+ txs worst case), against a close of one or two transactions. The contract reserve answers this by making the party who loses pay. It *does* imply, however, that a small-balance user might have to lock up a non trivial percentage in the reserve; this argues for usage at slightly higher balances (Section 10).

10. Costs

item	measured (proof of concept)
proving the withdrawal (STARK)	446k cycles with RISC Zero's accelerators, one segment: 2 min 24 s on a 12-core x86 laptop (4.8 GB peak), 30 s on an Apple-silicon Mac (M4 Pro, CPU); without the accelerators, 5.9M cycles, almost all the sequencer's signature
the Groth16 wrap	109 s (x86 only)
the verifier on the 172-byte input	478,100,841 steps; 58 s in BitVMX's emulator
the same proof checked natively (the pairing, RISC Zero's claim digest)	1.5 ms
BitVMX's 8-ary search over it, a jump list (both parties' executions)	about 12 min
`escalate` (move 1: the claim, round 1's hashes, 11 signed input words)	11.6 kvB
`inputs` (the other 32 input words)	8.0 kvB
Alice's round (seven hashes), the hub's choice (each besides its poster's 200 vB child)	7.5 kvB, 0.3 kvB
the final record, the last step proved	6.8 kvB, 25.1 kvB
a read: `path_split`, `combine_`k (k chunks), the verdict	0.6 kvB, about 11.6 kvB a chunk, 23.6 kvB
the dispute's timeout, a bit output, an input-word check, `equiv_b`, `halt_exit`, the default	430 vB, ∼300 vB, 413 vB, 640 vB, 5.1 kvB, 280 vB

Terminal transactions include the broadcaster's fee input (about 60 vB).

A dispute. A refused withdrawal puts the whole search on chain. Measured on regtest under mainnet's relay policy, a valid claim played through and paid is 123 kvB in 22 transactions, most of it Alice's rounds of hashes (92 kvB) and the last step's proof (25 kvB); a false claim disproved is 103 kvB. A lie about a memory read adds the jump list's path: in the case measured, 25 transactions and about 187 kvB in all, with a worst-case path's combine at 93 kvB. These bytes, at whatever feerate prevails, are what the loser's reserve pays for. lichen's hub and wallet played these disputes against each other on the Groth16 verifier, a valid claim proved in 725 s and a false one disproved in 540 s on regtest; an honest hub that checks the proof natively concedes a true claim without playing at all. On mainnet the time is set by the move windows: with each side using its whole window, a dispute takes about $11(w+1)$ blocks, some 140 at $w = 12$. Further reductions:

- *a shorter jump list*: Lerner's tighter long-jump function (about $\log_2 n + 3\log_2\log_2 n$ hops, some 45 instead of 86), or a counter that moves over a step number's set bits only, would shrink a read's combine about in proportion;
- *a chess clock*: each move's timeout waits a fixed w blocks, so a party that stalls on every move stretches the dispute to the $11(w+1)$ blocks above. Union [11] instead gives each

party a budget for all its moves together: while one party is due to move, the other publishes timelocked transactions marking the blocks that pass, and the mover's next transaction commits to the time it used under a one-time signature, so a timeout fires when a party's total exceeds its budget, not when one wait does. The worst case falls to the two budgets plus about a block a move; the price is the extra transactions, pre-signed, and published when a party stalls.

Can a user dispute on her own hardware? A dispute needs two computations from Alice, neither of them urgent (only T_{close} bounds a withdrawal), and none at all in the cooperative case, where the hub reads the layer 2 directly. First, a proof of her withdrawal. As built, the statement costs 446,000 zkVM cycles, one segment, using RISC Zero's accelerators (about 10 cycles for a 256-bit modular multiplication, 68 for a SHA-256 block [16]); without them it was 5.9 million, almost all of them in the sequencer's signature. Proving it takes 30 s on an Apple-silicon laptop's CPU and 2 min 24 s on a 12-core x86 laptop, with 4.8 GB of memory at peak, which smaller segments reduce [17]; a CUDA GPU would be faster still. The Groth16 wrap is a fixed cost of a few minutes, today on x86 only; it can be delegated, since the wrapper sees only the receipt and its public journal. Second, her side of the search: one run of the verifier with the trace hashed, a few minutes of CPU on a laptop. That run comes before her first move, since the escalation commits to the whole trace, so the escalation follows a refused fold by those minutes at least; against windows of hours, and a withdrawal bounded only by T_{close}, that is no constraint. Both are within a consumer laptop's reach; a phone is plausible for the proof with small segments and a delegated wrap, and less so for the dispute onchain, as discussed in Watchtowers 7.

Pre-signing. About two transactions a move, the final step's proofs (62 instruction classes for the Groth16 verifier) and the read and jump challenges' set (the split, eight combines, their choices and timeouts): 141 per contract, measured; inside a channel, that set for each of the two commitments, re-signed on every channel update, if we choose to put contracts inside Lightning channels (which is probably the more user-convenient version).

11. Positioning

We have already briefly discussed how this is different from just "payments over Lightning via a hub" in the Introduction. We'll list here some key differences to some other existing L2 or L2-like Bitcoin constructions.

Against BitVM2 and Union. The closest relatives, technically, are two bridges: BitVM2 [10], and Union [11], a bridge to Rootstock built on BitVMX by BitVMX's own authors. All three release bitcoin on Bitcoin against a claim about a program's execution, settled by fraud proof, with no change to consensus, and all three inherit the assumptions of a side system they do not secure. They differ in who the parties are, and most of the rest follows (Table 3).

The bridges pool their users' bitcoin, and a pool shared by users not yet known must be released by transactions that someone signs in advance on their behalf: BitVM2's committee, Union's functionaries. What then shapes the dispute is who may challenge a false claim. BitVM2 opens challenge to anyone, since anyone might be the one robbed; with the challengers unknown when the graph is signed, no interactive game can be pre-signed with them, so its dispute is

	a BitVM2 bridge	Union	lichen
custody	one pool for all users	one pool, in fixed-size outputs, one per peg-in	a contract for each user, with the hub
pre-signed by	a committee, for every operator's graph on every deposit	the N functionaries, a two-party game for every pair of them on every peg-in	the user and the hub, 141 transactions per contract
setup trust	one honest committee member	one honest functionary	none
the claimant	an operator, who paid the user from its own funds and claims reimbursement	a functionary, likewise	the user, for her own withdrawal
the challengers	anyone; safety needs one honest one watching	the other functionaries; safety needs one honest one watching	the hub alone, the only party a false claim can rob
the dispute	one round: the operator asserts the verifier's intermediate values, a challenger reveals one failing chunk (scripts up to 4 MB); 42 blocks	BitVMX's interactive search, 4-ary: about 270 kvB at worst; a heavier chain is answered in a nested game	BitVMX's interactive search, 8-ary, with a jump list: 22 transactions, 123 kvB; a read in three or four more
the user's part in a dispute	none: she was paid up front	none, likewise	a party: she plays the search, or a watchtower does
a refused withdrawal	it waits for a willing operator	it waits for a willing functionary	she enforces it on chain herself
stakes	operators' collateral; challengers fund the assert	bonds sized to dispute costs, shared by a packet of peg-ins	the contract reserves r_A, r_H
liquidity	operators front withdrawals until reimbursed	functionaries likewise, one peg-out each at a time	the hub locks $V_{\max} - D + r_H$ for each user until T_{close}
the asset	bridged bitcoin	bridged bitcoin, in fixed denominations	the hub's tokens: bitcoin within each user's cap, credit beyond it
expiry	none	none	T_{close}: withdraw, or renew, before it

Table 3: Two bridges and lichen. The bridges' figures are their authors' estimates; lichen's are measured on regtest (Section 10).

a single round, in which the operator puts the verifier's whole intermediate state on chain at once. Union closes challenge to its functionaries, who are known at setup, and so can pre-sign BitVMX's interactive search between every pair of them, which puts less data on chain over many more rounds. Its price is that safety needs an honest functionary not only at setup but online whenever a claim is made, and that the pairwise games, and the bonds that fund them, grow with the number of functionaries. A lichen contract is not pooled: it has two parties, both present when it opens. They pre-sign it themselves, so there is no committee and no setup trust; only the hub can be robbed by a false claim, so only the hub challenges; and the search between them is Union's game with a single pair. Splitting the pool has its own price: the hub's locked liquidity, a token that is bitcoin only within each user's cap, an expiry, and a user who must take part in her own dispute. **The bridges ask their users to trust one signer at setup and one challenger afterwards, and to be present for none of it; lichen asks them to trust no one, but to be present for their own disputes**.

Other designs, briefly.

- *Ecash* (Cashu, Fedimint): lichen removes the mint's power to take the funds, within each user's cap, at the price of the hub's locked liquidity; the mint may print as it likes, but every redemption is against its own collateral.
- *Ark*: a shared tree of outputs with an operator and an expiry, where a lichen contract is per user; both ask users to act before an expiry; Ark's holdings are direct claims on bitcoin, lichen's are the hub's tokens, enforceable within the cap, on a programmable ledger.
- *Payment pools* (one shared output, balances updated by all members): need every member's signature for an update and a covenant (or an exponential tree) for unilateral exit; a lichen transfer needs only its sender, and exit needs no consensus change.

12. Open questions

- *The layer 2's canonicity and inclusion*: the sequencer (single, distributed, or Bitcoin itself as in Shielded CSV [19]) and forced inclusion; deliberately separate from trust in the hub; roots sealed on Bitcoin with light-client proofs (Section 6); data availability.
- *A dispute's size and duration*: a chess clock in place of per-move timeouts; a shorter jump list (Lerner's tighter long-jump function, about $\log_2 n + 3\log_2\log_2 n$ hops, or a counter over only a step number's set bits), which shrinks a read's path and its combine (Section 10).
- *Fees*: links cannot be pinned (only the poster can spend an unconfirmed one), but they pay a floor of 1 sat/vB and so, when mempools are full, rely on package relay of the link with its poster's child. Terminal transactions, signed `ALL|ANYONECANPAY`, can be pinned through any output spendable before confirmation (the hub's share of a payout, the default's outputs); a `1 OP_CSV` on every such output, as on the links and as in Lightning's anchor channels, would close this, and a party posting early in its window never races a timeout. Fee coins are added without change, so a party keeps coins split to size.
- *Watchtowers*: is the "turnkey" solution actually practical? (Section 7).

13. Implementation

The design is implemented in proof-of-concept form as two code bases, LN-GAP [3] and Lichen [4], with tests and demos executed on bitcoin regtest. Any bitcoin-related data in this paper come from the concrete executions there, including playing out disputes fully.

The code is not feature complete, in particular in regard to the L2 itself which is very much "minimum-viable": transparent notes, single sequencer with single signature, and withdrawal proofs are against the sequencer's signed state root with no transition validity proof yet included, and no sealing of state roots on Bitcoin either.

References

[1] waxwing. LN-GAP: Lightning governed by arbitrary programs. Version 1.0 (draft), September 2026. https://github.com/AdamISZ/ln-gap/blob/master/docs/paper/lngap_draft.pdf

[2] waxwing. From a Rust program to one step on Bitcoin: how a withdrawal is proved, wrapped and disputed in LN-GAP 2.5. October 2026. https://github.com/AdamISZ/ln-gap/blob/master/docs/paper/pipeline.pdf

[3] LN-GAP proof-of-concept code. https://github.com/AdamISZ/ln-gap

[4] lichen: the ledger, the sequencer, the hub and the wallet. https://github.com/AdamISZ/lichen

[5] S. D. Lerner, R. Amela, S. Mishra, M. Jonas and J. Álvarez Cid-Fuentes. BitVMX: a CPU for universal computation on Bitcoin, 2024. https://arxiv.org/abs/2405.06842

[6] FairgateLabs. BitVMX-CPU: the emulator, the n-ary search and the Groth16 verifier for BitVMX. https://github.com/FairgateLabs/BitVMX-CPU

[7] BitVMX-CPU with LN-GAP's changes (two bug fixes, the authenticated jump list), branch `lngap`. https://github.com/AdamISZ/BitVMX-CPU

[8] J. Poon and T. Dryja. The Bitcoin Lightning Network: scalable off-chain instant payments, 2016. https://lightning.network/lightning-network-paper.pdf

[9] R. Linus. BitVM: compute anything on Bitcoin, 2023. https://bitvm.org/bitvm.pdf

[10] R. Linus Woll, L. Aumayr, Z. Avarikioti, M. Maffei, A. Pelosi, O. S. Thyfronitis Litos, C. Stefo, D. Tse and A. Zamyatin. Bridging Bitcoin to second layers via BitVM2. USENIX Security 2026; ePrint 2025/1158. https://eprint.iacr.org/2025/1158

[11] R. Amela, S. Mishra, S. D. Lerner and J. Álvarez Cid-Fuentes. Union: a trust-minimized bridge for Rootstock, 2025. https://arxiv.org/abs/2501.07435

[12] RISC Zero. The RISC Zero zkVM and its Groth16 receipts. https://dev.risczero.com

[13] S. D. Lerner. Latest innovations in BitVMX (trace Merkle tree, authenticated skiplist and jumplist), June 2024. https://bitvmx.org/knowledge/latest-innovations-in-bitvmx

[14] J. Harris and A. Zohar. Flood & Loot: a systemic attack on the Lightning Network. AFT 2020. https://arxiv.org/abs/2006.08513

[15] P. Todd. Preventing consensus fraud with commitments and single-use-seals, 2016. https://petertodd.org/2016/commitments-and-single-use-seals

[16] RISC Zero. Cryptography acceleration (SHA-256 and 256-bit modular multiplication circuits; patched `k256`, `sha2`). https://dev.risczero.com/api/zkvm/acceleration

[17] RISC Zero. zkVM datasheet: proving time and memory by cycle count, CPU, Metal and CUDA, April 2023. https://dev.risczero.com/datasheet.pdf

[18] DLC specifications: numeric outcome compression. https://github.com/discreetlogcontracts/dlcspecs

[19] J. Nick, L. Eagen and R. Linus. Shielded CSV: private and efficient client-side validation. ePrint 2025/068. https://eprint.iacr.org/2025/068

[20] C. Shikhelman, M. Komarov and A. Moskvin. Shielded Bitcoin: private transfers on the Bitcoin L1. [alloc] init, September 2026. https://www.allocinit.xyz/uploads/shielded-bitcoin.pdf